

# Traditional Parallel Computing vs Parallel Computing with Cloud

---

Andrew Xie, Antonio Mena, Prayag Patel

# Background

---

# Project Goal & Implementation

**Goal**: Compare traditional parallel computing vs parallel computing with cloud


**Implementation Task**: Gravity simulator that tracks particle trajectories

## Why This Problem?

- Classic scientific computing problem
- $O(n^2)$  computational complexity
- Naturally distributable workload

## Technologies Used:

- Traditional: Spark (pySpark)
- Cloud: MPI on AWS EC2 nodes



# Gravity Simulation Background

Particles in the simulation have:

- Position
- Velocity
- Mass

Newton's Law of Universal Gravitation

Applications:

- Predict cosmic collisions
- Future particle positions and their impact on other nearby cosmic objects
- Compare with real world data to verify known objects and identify unknown objects

$$F = G \frac{m_1 m_2}{r^2}$$

$F$  = force

$G$  = gravitational constant

$m_1$  = mass of object 1

$m_2$  = mass of object 2

$r$  = distance between centers of the masses

# What is Amazon EC2?

- **Amazon Elastic Compute Cloud (EC2)** provides scalable, on-demand computing resources in the AWS Cloud.
- EC2 enables you to launch virtual servers (**instances**) and scale resources based on workload requirements, reducing hardware costs and enabling faster application development and deployment.



# How do EC2 Instances Work?

- **Instances:** Virtual servers that run applications and processes in the cloud. You choose the instance type based on your project's resource needs (CPU, memory, storage, and networking).
- **Instance Types:** Each type offers different combinations of resources to support different workloads, ideal for tasks like parallel computing with MPI.
- **Amazon Machine Images (AMIs):** Pre-configured templates for your instances, which include operating systems and additional software required for your project.
- **Instance Lifecycle:** You can start, stop, or terminate instances as needed. When an instance is terminated, its resources (such as temporary storage) are released.

# Related Work

—

# Existing Comparisons Between Spark and MPI

- Using the Cloud for parameter estimation problems: comparing Spark vs MPI with a case-study (*Gonzalez, Pardo, et al.*) [\[1\]](#)
- Performance Evaluation of Apache Spark Vs MPI: A Practical Case Study on Twitter Sentiment Analysis (*Kumar, Rahman*) [\[2\]](#)
- Big Data in metagenomics: Apache Spark vs MPI (*Abuín, Lopes, et al.*) [\[3\]](#)
- Comparing Spark vs MPI/OpenMP On Word Count MapReduce (*Junhao Li*) [\[4\]](#)

Existing studies similarly evaluate Spark and MPI frameworks against different use cases.

This project bridges the gap between scientific and big data workloads:

- ★ gravity simulation employs both **compute-heavy calculations** and **complex communication** for particle interactions

This project emphasizes parallel **cloud** computing:

- ★ our MPI experiment is run directly on **AWS EC2** cloud infrastructure, reflecting real-world cloud computing scenarios



# Traditional Parallel Computing (Spark)

—

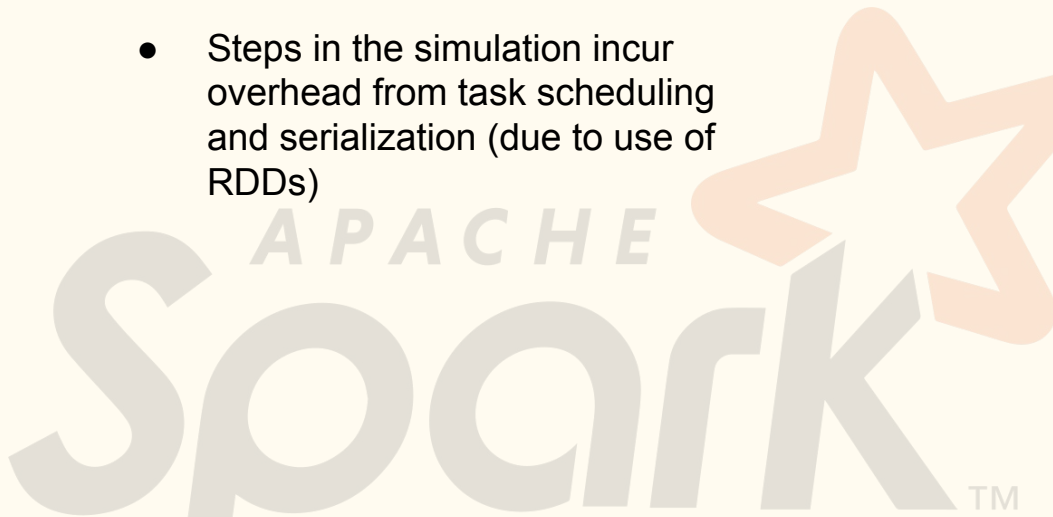
# Traditional Parallel Computing (Spark)

## Implementation details:

- Leveraged `pySpark` with `numpy` and `dataclasses`
- `Particle` class: stores 3D position, velocity, mass
- Force calculations using broadcasted particle data
- Algorithm steps:
  - 1) Force recalculation
  - 2) State updates
  - 3) Trajectory recording

## Key Challenges:

- Potential bottleneck from data shuffle costs
- Not natively optimized for iterative workloads
- Steps in the simulation incur overhead from task scheduling and serialization (due to use of RDDs)




# Parallel Cloud Computing (MPI on EC2 nodes)

—

# Parallel Cloud Computing (MPI on EC2 Nodes)

## Implementation approach:

- Custom `Particle` data structure
- Vector utility functions for 3D arithmetic
- Even distribution of work across MPI “ranks”
- Global synchronization using `MPI_Allgatherv`



MPI built-in that gathers data from all tasks and broadcasts the data to all tasks

## AWS setup:

- Used `t2.micro` instances (free tier)
- 1 master node + 2 slave nodes
- Passwordless SSH for node communication
- Horizontal scaling to overcome single-processor limitation
- Leverage EC2’s placement group to colocate instances and reduce communication latency

EC2

# Initial Parameters

```
def create_solar_system() -> List[Particle]:  
    return [  
        Particle(  
            position=np.array([0.0, 0.0, 0.0]),  
            velocity=np.array([0.0, 0.0, 0.0]),  
            mass=1.989e30  
        ),  
        Particle(  
            position=np.array([1.496e11, 0.0, 0.0]),  
            velocity=np.array([0.0, 29.78e3, 0.0]),  
            mass=5.972e24  
        ),  
        Particle(  
            position=np.array([2.279e11, 0.0, 0.0]),  
            velocity=np.array([0.0, 24.077e3, 0.0]),  
            mass=6.39e23  
        )  
    ]
```

```
particles[0].position = vector3_zero();  
particles[0].velocity = vector3_zero();  
particles[0].mass = 1.989e30;
```

```
particles[1].position.x = 1.496e11;  
particles[1].position.y = 0.0;  
particles[1].position.z = 0.0;  
particles[1].velocity.x = 0.0;  
particles[1].velocity.y = 29.78e3;  
particles[1].velocity.z = 0.0;  
particles[1].mass = 5.972e24;
```

```
particles[2].position.x = 2.279e11;  
particles[2].position.y = 0.0;  
particles[2].position.z = 0.0;  
particles[2].velocity.x = 0.0;  
particles[2].velocity.y = 24.077e3;  
particles[2].velocity.z = 0.0;  
particles[2].mass = 6.39e23;
```

# Random Particles

```
for i in range(5):  
    more_particles.append(Particle(  
        position=np.random.uniform(-3e11, 3e11, 3),  
        velocity=np.random.uniform(-30e3, 30e3, 3),  
        mass=np.random.uniform(1e23, 1e25)  
    ))
```

```
srand(time(NULL));  
for (int i = 3; i < 8; i++) {  
    particles[i].position.x = ((double)rand() / RAND_MAX) * 6e11 - 3e11;  
    particles[i].position.y = ((double)rand() / RAND_MAX) * 6e11 - 3e11;  
    particles[i].position.z = ((double)rand() / RAND_MAX) * 6e11 - 3e11;  
    particles[i].velocity.x = ((double)rand() / RAND_MAX) * 6e4 - 3e4;  
    particles[i].velocity.y = ((double)rand() / RAND_MAX) * 6e4 - 3e4;  
    particles[i].velocity.z = ((double)rand() / RAND_MAX) * 6e4 - 3e4;  
    particles[i].mass = ((double)rand() / RAND_MAX) * 9.9e24 + 1e23;  
}
```

# Force Calculation

```
def calculate_force_between(p1_data, p2_data, G):  
    p1_pos = np.array(p1_data['position'])  
    p2_pos = np.array(p2_data['position'])  
    r = p2_pos - p1_pos  
    distance = np.linalg.norm(r)  
  
    if distance < 1e-10:  
        return np.zeros(3).tolist()  
  
    force_magnitude = (G * p1_data['mass'] * p2_data['mass']) / (distance ** 2)  
    return (force_magnitude * r / distance).tolist()
```

```
Vector3 calculate_force(Particle* p1, Particle* p2) {  
    Vector3 force = vector3_zero();  
    Vector3 r = vector3_subtract(p2->position, p1->position);  
    double distance = vector3_magnitude(r);  
  
    if (distance < 1e-10) {  
        return force;  
    }  
  
    double force_magnitude = (G * p1->mass * p2->mass) / (distance * distance);  
    double scale = force_magnitude / distance;  
    force = vector3_multiply(r, scale);  
  
    return force;  
}
```

# Parallelization

```
class SparkGravitySimulator:
    G = 6.67430e-11

    def __init__(self, particles: List[Particle], dt: float = 0.01):
        self.spark = SparkSession.builder \
            .appName("GravitySimulation") \
            .config("spark.executor.memory", "2g") \
            .getOrCreate()

        self.particles_data = [p.to_dict() for p in particles]
        self.dt = dt
        self.num_particles = len(particles)

    def calculate_forces(self):
        particle_pairs = []
        for i in range(self.num_particles):
            for j in range(i + 1, self.num_particles):
                particle_pairs.append((i, j))

        sc = self.spark.sparkContext
        particles_broadcast = sc.broadcast(self.particles_data)
        G_broadcast = sc.broadcast(self.G)

        pairs_rdd = sc.parallelize(particle_pairs)

        def calculate_pair_force(pair):
            i, j = pair
            particles = particles_broadcast.value
            G = G_broadcast.value
            force = calculate_force_between(particles[i], particles[j], G)
            return (i, j, force)

        forces = pairs_rdd.map(calculate_pair_force).collect()
```

```
MPI_Get_address(&all_particles[0].position, &displacements[0]);
MPI_Get_address(&all_particles[0].velocity, &displacements[1]);
MPI_Get_address(&all_particles[0].mass, &displacements[2]);

for (int i = 2; i <= 0; i--) {
    displacements[i] = MPI_Aint_diff(displacements[i], displacements[0]);
}

MPI_Type_create_struct(3, blocklengths, displacements, types, &particle_type);
MPI_Type_commit(&particle_type);

MPI_Bcast(all_particles, num_particles, particle_type, 0, MPI_COMM_WORLD);

MPI_Allgatherv(
    &all_particles[start_idx], local_num_particles, particle_type,
    all_particles, recvcunts, displs, particle_type,
    MPI_COMM_WORLD
);

free(recvcunts);
free(displs);

MPI_Barrier(MPI_COMM_WORLD);
```



# Output

Starting gravity simulation at 20241211\_020916

Configuration:

- Number of steps: 500
- Time step: 3600 seconds (1 hour)

Performance Statistics:

Total execution time: 55.43 seconds

Average time per step: 0.1109 seconds

Final positions:

Particle 0: (-40759.92361748901, 83021.92422931321,  
-15339.827766454993)  
Particle 1: (140072929991.45477, 52463379787.609665,  
-21994.541876042145)  
Particle 2: (223763210109.76465, 43076578937.28692,  
-20591.140443501015)  
Particle 3: (-220536191887.73007, -131004855333.84792,  
-270275726748.04044)  
Particle 4: (69353928904.48807, 203022216340.21124,  
-168798879233.916)  
Particle 5: (-203270117941.5046, -190823886331.11755,  
-272352403694.52585)  
Particle 6: (216587611342.19537, -42911503641.20782,  
-259177568462.8093)  
Particle 7: (-26352676432.34539, 13632601778.28655,  
23526204261.72007)

Simulation completed successfully

Starting MPI C gravity simulation at 20241211\_015031

Number of processes: 3

Number of particles: 8

Steps: 500

Timestep: 3600.000000 seconds

Performance Statistics:

Total execution time: 0.28 seconds

Average time per step: 0.0006 seconds

Final positions:


Particle 0: (2.038543e+04, 7.519318e+03, 4.330912e+03)  
Particle 1: (1.400729e+11, 5.246338e+10, 4.085081e+03)  
Particle 2: (2.237632e+11, 4.307658e+10, 3.633190e+03)  
Particle 3: (3.104993e+11, 7.523407e+10, 2.527527e+11)  
Particle 4: (-7.151560e+09, 1.482256e+11, 1.320689e+11)  
Particle 5: (-1.181584e+11, 1.983205e+11, 2.194526e+11)  
Particle 6: (-2.203978e+11, -6.206211e+10,  
-1.531738e+11)  
Particle 7: (-1.774408e+11, -1.023138e+11,  
3.362449e+10)

Simulation completed successfully

# Results & Conclusion

—

# Results



	<b>Traditional Parallel Computing (Spark)</b>	<b>Parallel Cloud Computing (MPI on EC2 nodes)</b>
Total execution time	55.43 seconds	0.28 seconds
Average time per step	0.1109 seconds	0.0006 seconds



**~196.43 times speedup**


# PySpark vs MPI: Differences

- ★ MPI obviously performs better than PySpark because the gravity simulator used on AWS is written in C (statically typed compiled language vs dynamic interpreted Python)

**MPI is low-level** and requires better coding practices, it does not handle fault tolerance and must be written in C



**PySpark is high-level**, relatively easier to implement and has a higher level of abstraction



# Limitations of Spark & Traditional Parallel Computing

As particle count is increased, the Spark implementation tends to **struggle more with growing communication overhead**, while MPI maintains relatively constant performance.

Spark is optimized for **data-parallel workloads** with minimal inter-task communication.

- Frequent communication between particles (tasks) in this problem is critical.
- At each simulation step, Spark needs to refresh data using RDDs, which implicitly incurs serialization, deserialization, and disk/memory storage overhead.

# PySpark vs MPI: Increasing Efficiency?

- However, some improvements can be made to the PySpark code by optimizing memory used by the RDD, vectorizing operations, etc.
  - Better use of broadcast variables and smarter partitioning could improve performance
- The C code could also be improved by optimizing MPI, using more sophisticated algorithms, data management, etc.
  - Implementation of non-blocking communication could improve performance

# Conclusion

## **Key insights:**

- MPI & EC2 implementation significantly outperformed Spark for the gravity simulation problem
- Spark has clear limitations for tightly-coupled, iterative problems
- Cloud computing (via Amazon EC2) enabled efficient scaling

## **Future considerations:**

- Optimization opportunities for both platforms
- Potential for hybrid approaches
- Further benchmarking and testing needed
- Exploration and comparison of approaches for different use cases and problems